

Open weights: the risk sits on both sides of the argument

Twenty-five firms, including NVIDIA, Meta, Microsoft, IBM, and Hugging Face, published a joint case for open-weight models on July 24, 2026. It is an advocacy document, and the signatories hold commercial positions in the outcome. It is also the clearest statement yet of a real tradeoff: the risks of releasing model weights and the risks of not releasing them are different in kind, and neither disappears by choosing the other.

RISKS OF RELEASING WEIGHTS

IRREVERSIBILITY

Release cannot be undone

Once weights are published they are outside the developer's control. Modified versions are difficult to trace and impossible to recall. There is no rollback, no patch, and no kill switch.

SAFEGUARDS

Guardrails are removable

Safety behavior trained into a model can be fine-tuned out of a downloaded copy at modest cost. Controls that depend on the provider enforcing them do not survive distribution.

ACCOUNTABILITY

The control owner goes missing

After several rounds of fine-tuning and redistribution, no single party holds both the knowledge of what changed and the obligation to answer for it. Residual risk lands on the deployer by default.

RISKS OF NOT RELEASING THEM

CONCENTRATION

Few points, single failures

Capability held behind a small number of closed providers creates concentrated operational and supply dependencies. A breach, an outage, or a policy change at one provider propagates across every firm relying on it.

OPACITY

Closed is not the same as safe

Closed models can be breached, misused, or fail in ways outsiders cannot detect. External researchers cannot red-team what they cannot run, so verification depends on the provider's own account of itself.

DEFENSE

Asymmetric cyber capability

If attackers use advanced models, defenders need comparable capability to simulate and detect those threats. Restricting open weights constrains the defensive side more reliably than the offensive one.

Three questions the letter raises but does not settle

01

Where does the evidence trail start?

A deployer of a twice-fine-tuned open model needs provenance for weights, data, and modifications. Today that record usually begins at download, not at training.

02

What separates distillation from extraction?

The letter treats training on another model's outputs as legitimate practice and unlawful value extraction as distinct. The line between them is asserted, not drawn.

03

Who attests to an open model in production?

Closed providers can be contracted for assurances. Open weights carry no counterparty, so the attestation obligation shifts entirely to the deploying firm.

What this changes for a risk function

CONTROL AREA	THE OPEN-WEIGHT VERSION OF THE PROBLEM	PRACTICAL FIRST MOVE
Model inventory	Weights can be copied into an environment without a procurement event to trigger registration.	Detect at the infrastructure layer, not the vendor layer.
Third-party risk	There is no supplier to assess, no contract, and no service level to enforce.	Treat the model as an internally owned asset from day one.
Change management	Fine-tuning alters behavior without changing anything a release process would notice.	Version weights and evaluations together, or neither is meaningful.
Concentration risk	Closed-only strategies concentrate dependency; open-only strategies concentrate accountability.	Set a deliberate mix with named thresholds rather than defaulting to one.

Source: "Open Weights and American AI Leadership," joint open letter, July 24, 2026, published by NVIDIA and signed by 25 organizations including Andreessen Horowitz, Box, CrowdStrike, Dell Technologies, Hugging Face, IBM, The Linux Foundation, Meta, Microsoft, Mistral, Mozilla, Palantir, Perplexity, Replit, ServiceNow, and Y Combinator. Left column and right column risks are drawn from the letter's own text; the questions and control-area table are the author's analysis and are not claims made by the signatories.

VIEWS EXPRESSED ARE MY OWN AND DO NOT REPRESENT THOSE OF MY EMPLOYER. NOTHING HERE IS INVESTMENT, LEGAL, OR REGULATORY ADVICE.

Hitesh Patel • Enterprise risk & AI governance Leader